

Education

Georgia Institute of Technology

Atlanta, Georgia

Master of Science, Computer Science with specialization in Machine Learning; GPA 3.8/4

May 2021

- MS Advisor: [Prof. Devi Parikh](#)

- Conducting research in Language Grounding - Vision and Language Navigation, Image Captioning
- Selected Coursework: Deep Learning, Deep Learning for Text, Reinforcement Learning, Machine Learning, Machine Learning for Trading, Graduate Algorithms

Indian Institute of Technology Bombay

Mumbai, India

Bachelor of Technology, Computer Science and Engineering

May 2016

- IIT-JEE 2012 **All India Rank 14** out of 500k candidates
- **All India Rank 8** in Nationwide Education and Scholarship Test 2014. Received scholarship for the same
- Selected Coursework: Computer Graphics, Operating Systems, Compilers, Digital Geometry Processing

Publications

Error-driven Pruning of Language Models for Virtual Assistants

ICASSP 2021

Sashank Gondala*, Lyan Verwimp*, Ernest Pusateri, Manos Tsagkias, Christophe Van Gysel

Abstract: Language models (LMs) for virtual assistants (VAs) are typically trained on large amounts of data, resulting in prohibitively large models which require excessive memory and/or cannot be used to serve user requests in real-time. Entropy pruning results in smaller models but with significant degradation of effectiveness in the tail of the user request distribution. We customize entropy pruning by allowing for a keep list of infrequent n-grams that require a more relaxed pruning threshold, and propose three methods to construct the keep list. Each method has its own advantages and disadvantages with respect to LM size, ASR accuracy and cost of constructing the keep list. Our best LM gives 8% average Word Error Rate (WER) reduction on a targeted test set, but is 3 times larger than the baseline. To reduce the size of the keep list and the resulting LM, we propose two discriminative methods to select a subset of n-grams, one based on recognition errors of synthesized audio and one based on a text-only model that approximates the recognition errors. We show that the approximate version, while being much cheaper than the 'real' version, can retain the majority of the observed WER reductions.

Work Experience

Apple (Cupertino, CA)

Jun 2020 - Aug 2020

Machine Learning Research Internship (AI/ML)

- Worked as a research intern in the Language Modeling team
- Explored ways to prune language model of a speech recognition system without affecting its accuracy
- Obtained **10% reduction in LM size** with negligible increase in WER
- Work published at ICASSP 2021

Decentralized CDN Startup (San Francisco, CA)

Feb 2018 - Dec 2018

Co-founder

- **Co-founded a startup** to provide decentralized CDN services by sharing the spare bandwidth and hard drive space of Internet users over blockchain.
- Worked on various aspects of startup ranging from hiring and meeting investors to writing technical whitepaper and product development
- IP developed include a prototype to support HLS video streams and a [whitepaper](#) that describes the challenges and solutions based on the SOTA techniques including Service Certificates, Probabilistic Micropayments, etc.
- The startup was eventually discontinued due to a lack of product-market fit.

Oracle HQ (Redwood City, CA)

Jul 2016 - Aug 2019

Senior Member of Technical Staff

- **Improved sorting time** of a C++ in-memory query engine **by 15%** by identifying bottlenecks and enhancing the code to use compile time code generation techniques (C++11 Variadic templates).
- **Improved query run time** of benchmark set **by 20%** by enhancing caching algorithm logic modifying cache seed logic to cache the data post relevant processing rather than raw data.
- Introduced a new query syntax to enable auto discovery of backend tables bypassing the current requirement of manual import. **Reduced each ongoing release time by a few weeks.** Used YACC, LEX, and C++.
- Improved cache hit rate by changing the internal load balancer logic to create a deterministic server-user mapping instead of a session based allocation.

Software Developer Internship

- Worked with [Amazon Custom](#), the team that deals with customized products
- Built an API test suite for services of Amazon Custom using TestNG in Java.

Research Projects**Vision and Language Navigation****Sep 2020 - Current**

- Building an agent for the 'Vision-and-Language Navigation in Continuous Environment' ([VLN-CE](#)) task.
- The task is to train an agent to follow navigation instructions in a simulated house. The agent gets only the first-person view of the environment and navigates using low-level actions
- Building a hierarchical planner + controller architecture - planner predicts an intermediate waypoint and controller navigates to the waypoint. This improves sample complexity and overcomes problems with long-range planning
- Incorporating semantic + free space maps for better grounding of waypoint predictions

Image Captioning without reference captions**Jan 2020 - May 2020**

- Led a project at [Prof. Devi Parikh](#) and [Prof. Dhruv Batra](#)'s lab to generate captions for images containing novel objects (i.e., objects without paired training data) using non-paired data
- Modeled a [CIDEr predictor](#) using a pretrained multi-modal transformer to predict the CIDEr score without needing access to reference captions
- Using VIFIDEL, SLOR, and predicted CIDEr values as rewards, trained an image captioning model to optimize for these metrics using policy-gradient methods (PPO)

Other Projects**Question Answering using Deep Learning****Oct 2019 - Dec 2019**

- Worked on various Question Answering tasks - Google's Natural Question Answering and Stanford's SQuAD 2.0
- Implemented approaches such as LSTM based co-attention models, augmented BERT models, Ensembles, etc.

Neural Machine Translation**Mar 2019 - Apr 2019**

- [Implemented](#) a sequence-to-sequence (Seq2Seq) network in PyTorch to translate Spanish text to English
- Used a Bidirectional LSTM with multiplicative attention as Encoder and a Unidirectional LSTM as Decoder

ML trading bot**Mar 2019 - Apr 2019**

- Created a random forests-based trading algorithm which takes in the stock price and market indicators to predict the movement of a stock

OpenAI Agents**Jan 2019 - Feb 2019**

- Trained agents to solve several of the OpenAI challenges, using a mix of Reinforcement Learning (RL) techniques such as Q-Learning, DQN, DDQN, and Policy Iteration as a part of Reinforcement Learning course

Scholastic Achievements

- Secured **All India Rank 14** in IIT-JEE out of 500k test-takers **2012**
- Secured **All India Rank 59** in EAMCET out of 300k test-takers **2012**
- Obtained **7th position** in State Mathematics Olympiad (APAMT) **2009**
- Was placed **National Top 1%** in several Astronomy, Physics, and Junior Science Olympiads **2010-12**
- Attended (**Top 35 students across India**) Indian National Astronomy Olympiad (INAO) and Indian National Junior Science Olympiad (INJSO) Orientation-cum-Selection Camps held by [HBCSE](#) **2010**

Teaching Assistantships

- CS 7643 - Deep Learning, Georgia Tech **Spring 2021**
- CS 7643 - Deep Learning, Georgia Tech **Fall 2020**
- CS 8803 - Systems for Machine Learning Research, Georgia Tech **Spring 2020**
- CS 7641 - Machine Learning, Georgia Tech **Fall 2019**
- CS 101 - Intro to Computer Programming, IIT Bombay **Spring 2016**

Technical Skills

- **Languages:** C++ (Expert) | Python (Expert) | Java (Intermediate) | Bash (Intermediate)
- **Others:** Scikit-learn (Expert) | PyTorch (Expert) | TensorFlow (Intermediate) | SQL (Intermediate)